

Moxin Li

Homepage: <https://li-moxin.github.io/HelloFromMoxin/>
Github: <https://github.com/li-moxin>

Mobile: +65 80381637 / +86 13683398638
Email: limoxin@u.nus.edu

Education

- National University of Singapore Aug. 2021 – Now
Ph.D. Candidate in Computer Science, supervised by Professor Tat-Seng Chua.
- Peking University Sep. 2016 – Jun. 2020
Bachelor in Intelligence Science and Technology, Yuanpei College.

Experience

- National University of Singapore Sep. 2020 – Jul. 2021
Research Intern, supervised by Professor Tat-Seng Chua and Professor Fuli Feng.
- Wangxuan Institute of Computer Technology, Peking University Jun. 2019 – Jun. 2020
Research Intern, supervised by Professor Rui Yan.

Research Interest

My research interest lies in trustworthy language models. I have focused on topics including:

- Robustness: robust question answering by counterfactual reasoning, robust prompt optimization
- Confidence calibration for mitigating hallucination
- Multi-objective alignment towards thorough and unbiased trustworthiness

I am also interested in LLM evaluation, LLM knowledge boundary, prompt optimization, and table-and-text QA.

Selected Publications & Preprints

16 papers published in top-tier conferences.

6 papers as (co)first author.

[Google Scholar](#) citation -> 207. h-index -> 9.

- [Moxin Li*](#), Yuantao Zhang*, Wenjie Wang, Wentao Shi, Zhuo Liu, Fuli Feng, Tat-Seng Chua. Self-improvement towards pareto optimality: Mitigating preference conflicts in multi-objective alignment. ACL (findings) 2025.
- [Moxin Li*](#), Yong Zhao*, Yang Deng, Wenxuan Zhang, Shuaiyi Li, Wenya Xie, See-Kiong Ng, Tat-Seng Chua. Knowledge Boundary of Large Language Models: A Survey. ACL 2025 (oral presentation).
- [Moxin Li](#), Wenjie Wang, Fuli Feng, Fengbin Zhu, Qifan Wang, Tat-Seng Chua. Think Twice Before Trusting: Self-Detection for Large Language Models through Comprehensive Answer Reflection. EMNLP (findings) 2024.
- [Moxin Li](#), Wenjie Wang, Fuli Feng, Yixin Cao, Jizhi Zhang, Tat-Seng Chua. Robust prompt optimization for large language models against distribution shifts. EMNLP 2023 (oral presentation).
- [Moxin Li](#), Wenjie Wang, Fuli Feng, Hanwang Zhang, Qifan Wang, Tat-Seng Chua. Hypothetical training for robust machine reading comprehension of tabular context. ACL (findings) 2023.
- [Moxin Li](#), Fuli Feng, Hanwang Zhang, Xiangnan He, Fengbin Zhu, Tat-Seng Chua. Learning to imagine: Integrating counterfactual thinking in neural discrete reasoning. ACL 2022.

(* denotes Equal Contribution or Core Contributor.)

Professional Service and Presentation

- Tutorial Organizer: “Unveiling Knowledge Boundary of Large Language Models for Trustworthy Information Access” SIGIR 2025
- Conference and Journal Reviewer: ACL ARR, EMNLP ARR, ICLR, AAAI, KDD, CL
- Conference Volunteer: ICAIF 2025, EMNLP 2024, WWW 2024, WSDM 2023
- Session Chair and Presenter: The 2022, 2023, 2024 Singapore Symposium on Natural Language Processing
- Presenter: Singapore ACM SIGKDD Symposium 2024
- Test Designer and Organizer: CCIR Cup 2022

Honors and Awards

- NUS Research Achievement Award National University of Singapore, 2023